



# DataStax Enterprise Architecture


Aaron Regis & Richard Henderson

12th Jan 2018

# Agenda

1	Topology and Data Structure
2	Lab 1: Cassandra Access and Cassandra Stress
3	Request Handling
4	Lab 2: Hands-on DSE CQL

# Design Goals and Objectives



- Continuously Available
- Master Less
- Fully Distributed
- Shared-Nothing Architecture
- Build In Replication
- Linear Scalability
- Scale out

Architecture


# Apache Cassandra™ Architecture

- Cluster layer
  - Amazon DynamoDB paper
  - masterless architecture
- Data-store layer
  - Google Big Table paper
  - Columns/columns family




- All nodes are peers
  - Including seed nodes
  - No master
  - Discovery through gossip
- Built-in replication
  - Simplify your architecture!

## Cluster



# Token Ranges



**Token Range :  $2^{63} - 2^{63}$**

Example with **Replication Factor 3**

N3 will own data for tokens 1 – 30

Token Range : 1-10, owned by N1,N2,N3

Token Range : 11-20, owned by N2,N3,N4

Token Range : 21-30, owned by N3,N4,N5



- Primary key
  - Partition key
  - Clustering columns
- Partitioner
  - Generates unique hash from partition key
- Replication strategy
  - Token hash determines starting point
  - Determines replica placement

# Cassandra Query Language


```
CREATE KEYSPACE negib WITH replication =  
{'class': 'NetworkTopologyStrategy', 'DC1': '3'} ← Replication Factor  
AND durable_writes = true;
```

```
CREATE TABLE negib.sales (  
    name text,  
    time int,  
    item text,  
    price double,  
    PRIMARY KEY (name, time) ← Partition Key  
) WITH CLUSTERING ORDER BY (time DESC)
```

```
SELECT * FROM sales where name='Thomas' OR name='Günther' AND time >=20160401;
```


# Data Distribution

user_id1			
user_id2			
user_id3			
user_id4			
user_id5			



Token = hash of #partition  $\rightarrow$  #node  
Data is evenly distributed and clock wise replicated

# Data Distribution



Token = hash of #partition  $\rightarrow$  #node


Data is evenly distributed and clock wise replicated

# Lab 1 : Accessing the cluster

Read and write request handling


# Coordinator node

Incoming requests (read/write)  
Coordinator node handles the request



Every node can be coordinator → masterless

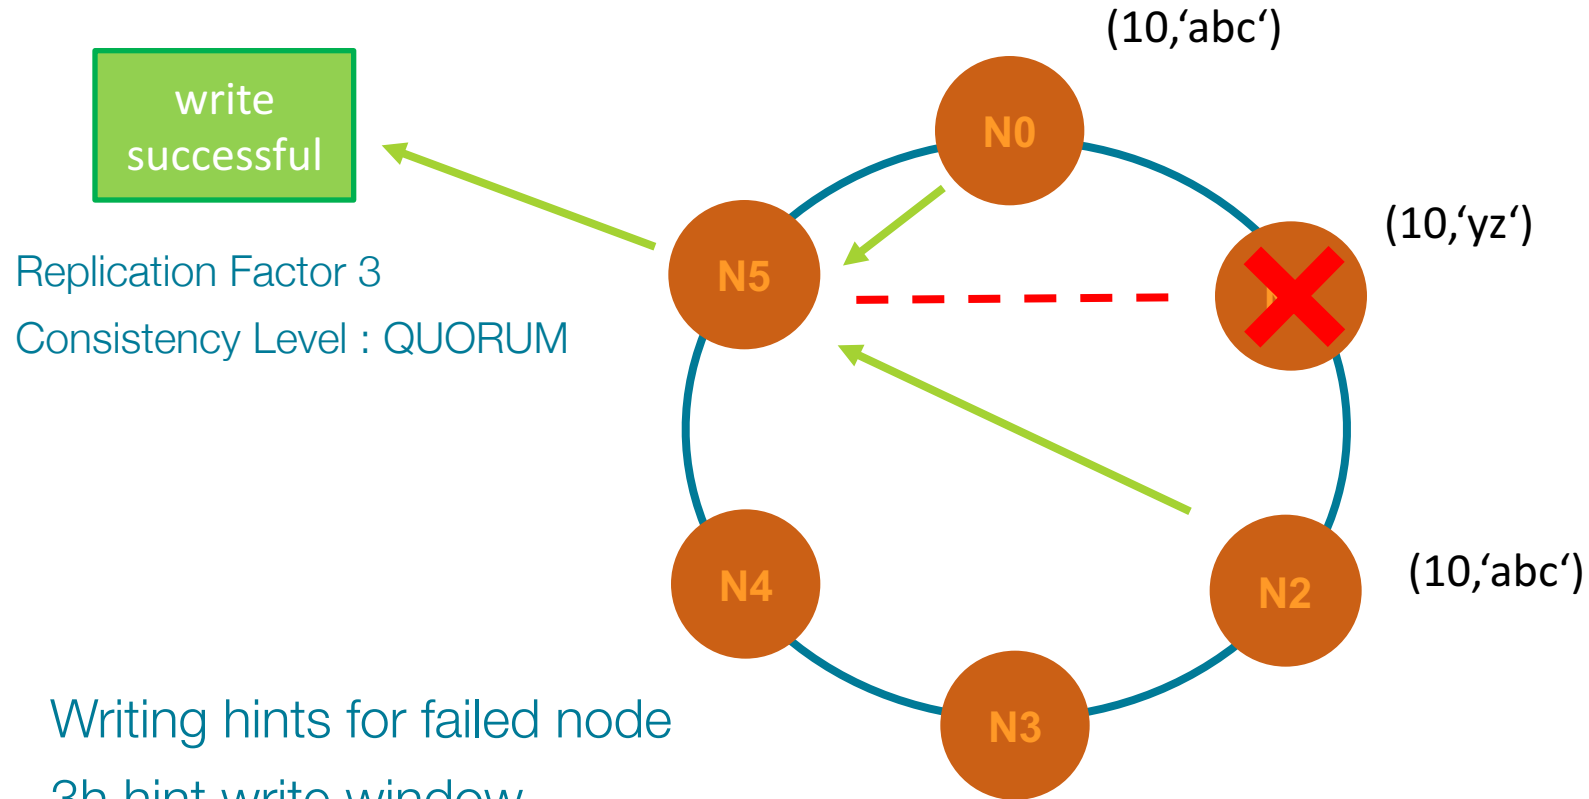
# Write request handling






- Background vs. Foreground Read Repair
  - Compare digests
    - If any mismatch
  - re-request to same nodes (full data set)
    - compare full data sets, send update
    - block until out-of-date replicas respond
  - Return merged data set to the client
- Consistency Level
  - one, quorum, all
  - local vs. cluster wide

# Write request handling



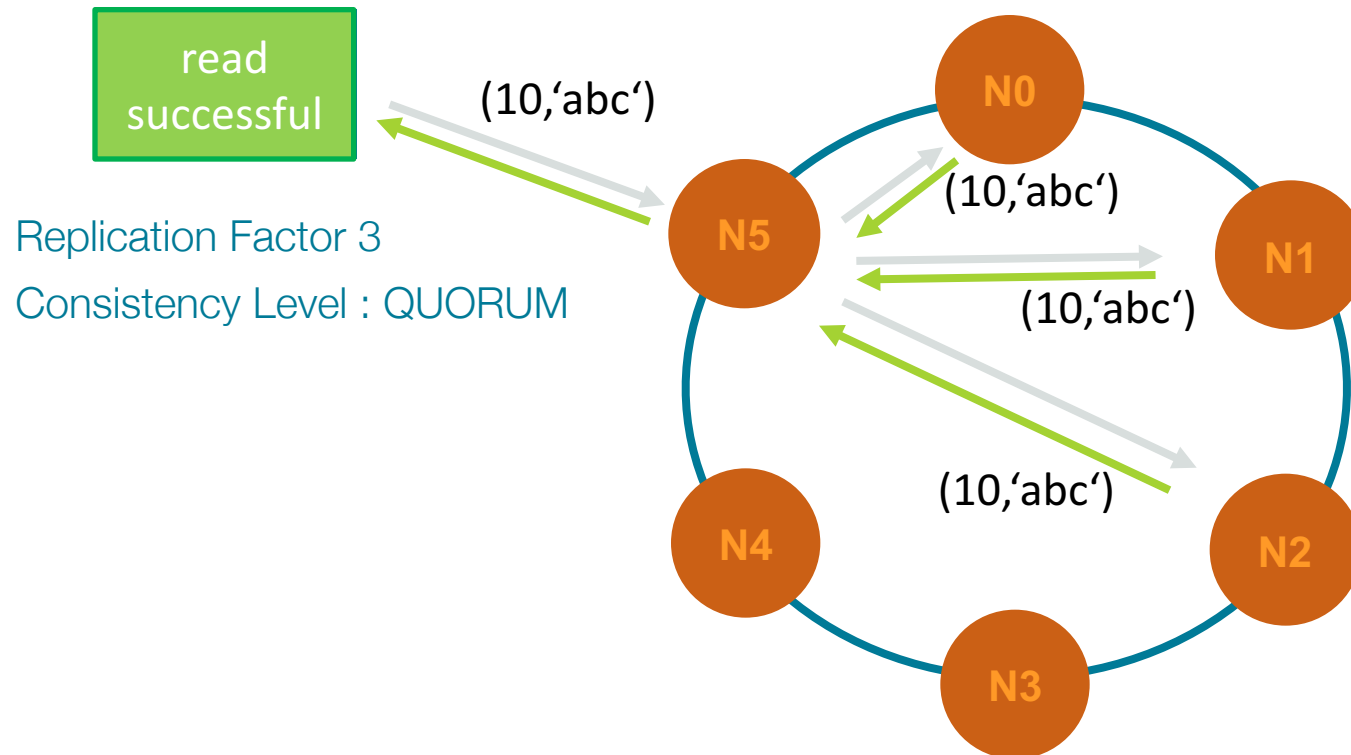
- Writing hints for failed node
- 3h hint write window
- CL: LOCAL\_ALL will fail

# Write request handling

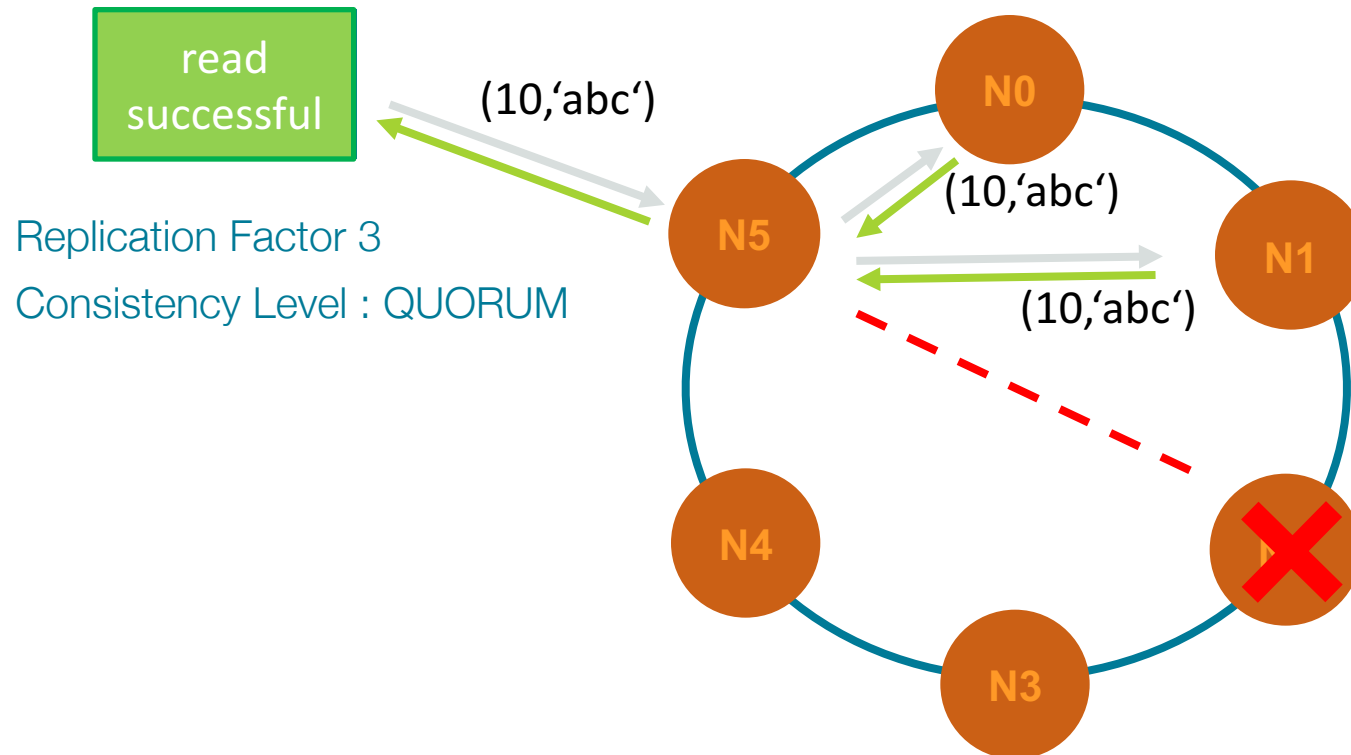


- quorum not met, failure
- CL: LOCAL\_ONE will succeed

# Read request handling

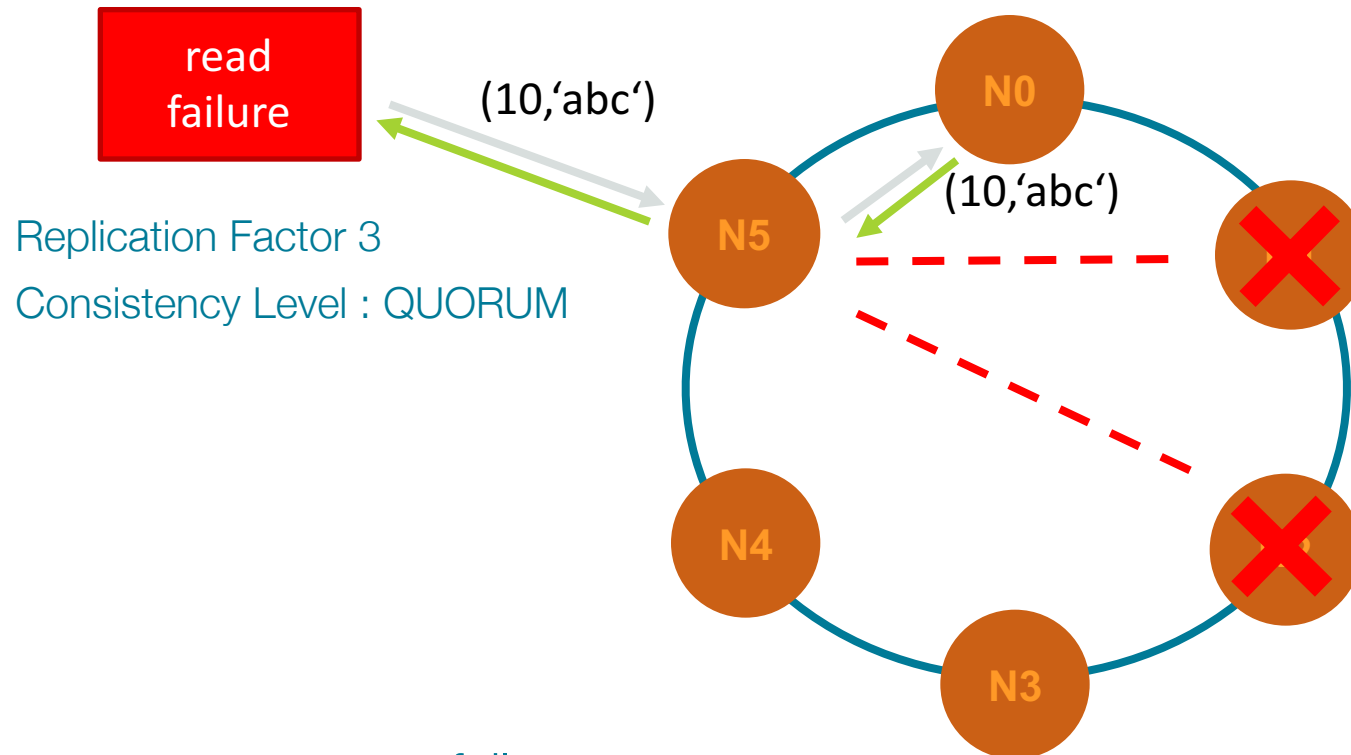


# Read request handling



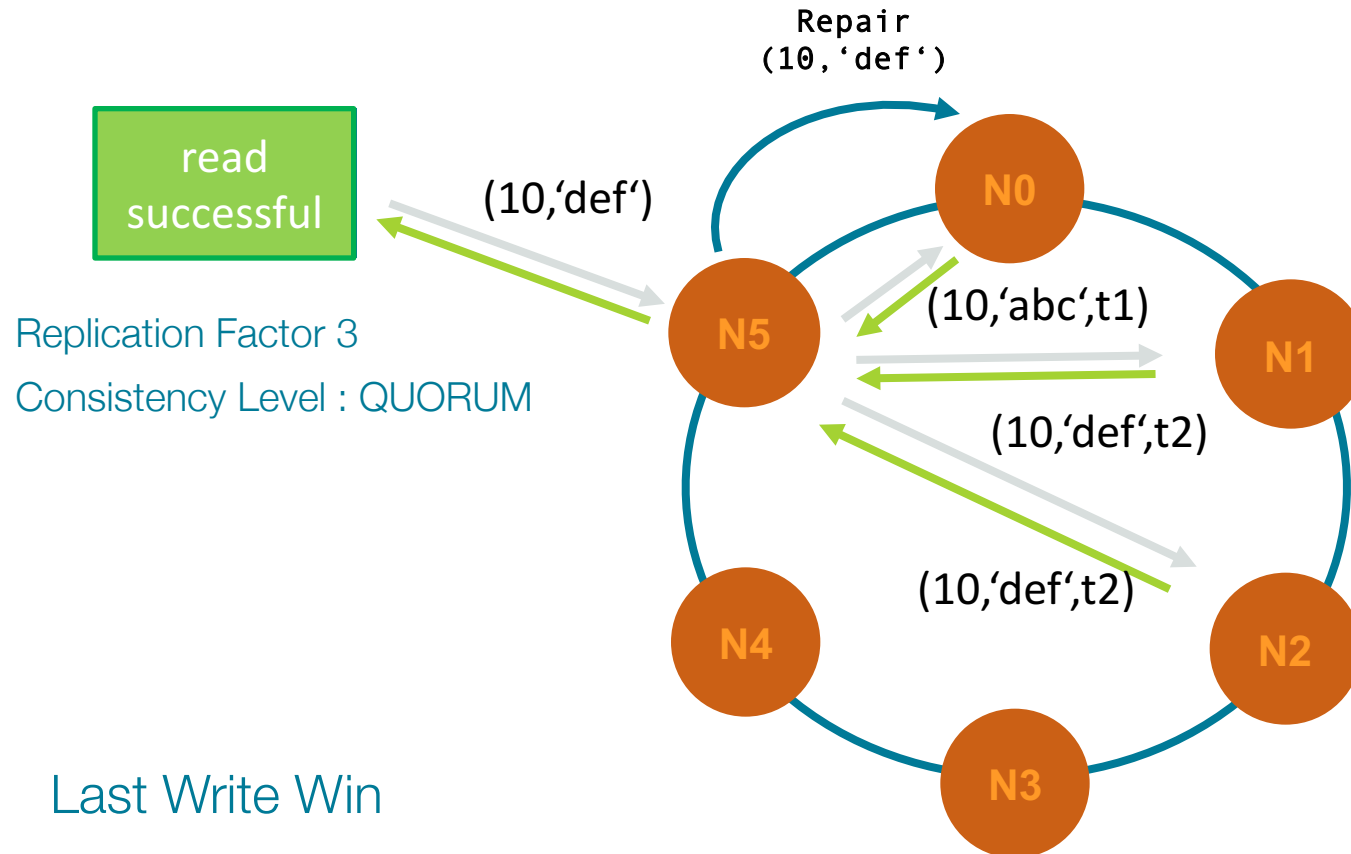
- Reading LOCAL\_QUORUM succeeds  
CL: LOCAL\_ALL will fail

# Read request handling



- quorum not met, failure
- CL: LOCAL\_ONE will succeed. [See more...](#)

# Read request handling – Read Repair



- Last Write Win
- $R+W > RF$  = immediate consistency
- Background vs. foreground Read Repair. [See more...](#)

# Driver Code

```
Cluster cluster = Cluster.builder()  
    .addContactPoint("127.0.0.1")  
    .withLoadBalancingPolicy(new TokenAwarePolicy(DCAwareRoundRobinPolicy.builder()  
        .withLocalDc("myLocalDC")  
        .build()))  
    .build();
```

```
PreparedStatement prepared = session.prepare  
    ( "insert into sales(id, name) values (?, ?)");
```

```
BoundStatement bound = prepared.bind("4", "Gregg");
```

```
session.execute(bound); // Throws UnavailableException If consistency doesn't met, downgrade is  
                           possible with corresponding RetryPolicy. Read More...
```



# Take away

- Data distribution (hash, tokens)
- Data replication (RF)
- All nodes are peer nodes , master less
- Background Read Repairs
- RetryPolicy in driver

# Lab 2 : Hands-on DSE CQL

Thank You!