



DataStax Enterprise Search

Rob Murphy & Aaron Regis


14th June 2017

Agenda

1

Introduction DSE Search

DSE Search




Live indexing engine with powerful search

- Automatic indexing on insert
- Higher ingestion throughput
- Distributed query optimization

Compared to self-managed:

- No separate search cluster to manage
- Probably less total hardware required
- No “Split Brain” data inconsistencies
- No ETL or synch to build and maintain
- No app level data management code

The data model is important



```
create table users (  
  user_id int,  
  name text,  
  age int,  
  gender boolean,  
  PRIMARY KEY(user_id) )
```

```
SELECT name FROM users WHERE user_id=1
```

```
SELECT * FROM users WHERE name="Thomas"
```

```
SELECT * FROM users WHERE age>45 and name="Thomas"
```

Not possible out of the box without
secondary idx or further extra tables

No problem with DSE Search

What is the value of DSE Search?

Multi-criteria WHERE Constraints

- WHERE constraints with multiple columns
- No extra tables needed

Full Text Search

- **Wildcards** ? *, like or Lemmatisation
- **Faceting**, Slice and Dice

Geospatial queries

- Queries with coordinates and distances search

Integrated with CQL

1. Create core, schema.xml and solrconfig.xml.
2. [optional] customize schema.xml, solrconfig.xml
3. Start indexing, re-indexing

Wildcard search

```
cqlsh> select * from sales where solr_cql = ' {`fq`:`name=gre*`} ';
```

Facet query

```
cqlsh> SELECT * FROM sales WHERE solr_query='{ "q": "name:*", "facet": { "field": "item" } }';
```

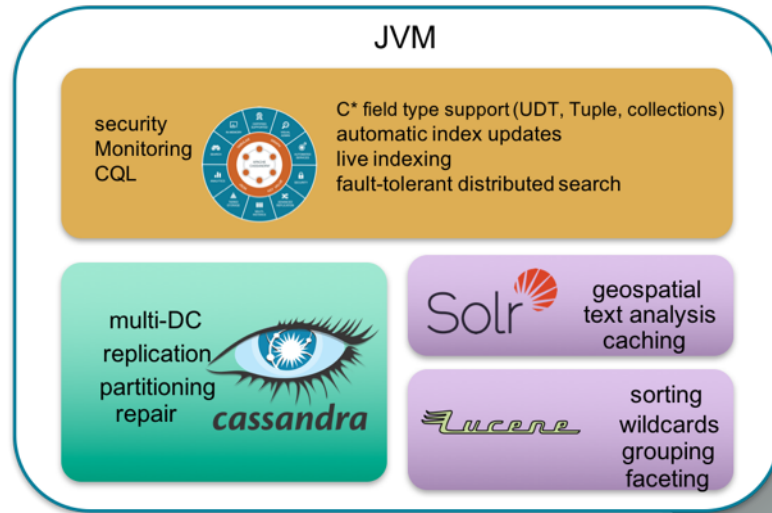
Range query

```
Cqlsh> SELECT * FROM sales WHERE solr_query='{ "q": "dt:[2017-01-01 TO 2017-01-10]" }';
```

Geo search

```
cqlsh> select * from sales where solr_cql = ' {"fq": "name=Hein*"} ';
```

DSE Search Architecture



Integrated in the the same JVM

- Data Locality (TokenRanges) and Shared Memory
- High Available, no master needed
- Index is sharded to all nodes

Transparency

- Access via CQL or REST API

Automated indexing with INSERT / UPDATE

- No extra ETL Process needed


With all benefits:

High Available


Scalable

Multi Data Center

Index Security



The Open Source Way



Separated Search Cluster


- “Split Brain” risk, data inconsistency
- ETL to generate, update and re-create index

Complex application

Two separated APIs, driver, read paths

https://docs.datastax.com/en/datastax_enterprise/5.0/datastax_enterprise/srch/searchOssSolrDiff.html

Inserting through CQL



Querying through CQL


```
SELECT title FROM solr WHERE solr_query='title:natio*';
```

title

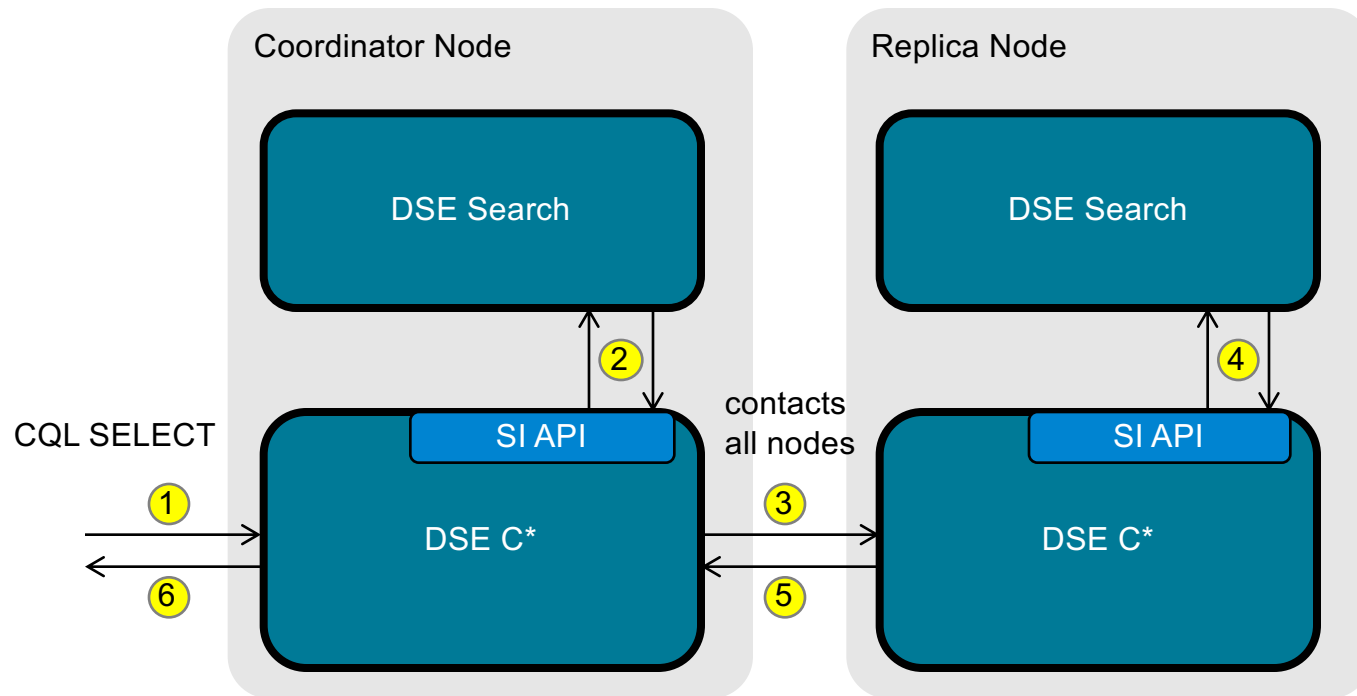
Bolivia national football team 2002
List of French born footballers who have played for other national teams
Lithuania national basketball team at Eurobasket 2009
Bolivia national football team 2000
Kenya national under-20 football team
Bolivia national football team 1999
Israel men's national inline hockey team
Bolivia national football team 2001

How many nodes to contact?

- We don't know the primary key
- Theory:
 - contact at least one replica for every token range
- Cassandra contacts all nodes
- **Our custom Solr SearchComponent does intelligent shard selection**



Querying through CQL



Lab 5 : Hands-on DSE Search

Thank You!