DataStax Enterprise Search

Rob Murphy & Aaron Regis

14th June 2017

Agenda

DATASTAX

DSE Search

Your Application

`CQL`

Search + Cassandra

Live indexing engine with powerful search
- Automatic indexing on insert
- Higher ingestion throughput
- Distributed query optimization

Compared to self-managed:
- No separate search cluster to manage
- Probably less total hardware required
- No "Split Brain" data inconsistencies
- No ETL or synch to build and maintain
- No app level data management code

DATASTAX

The data model is important

Partition Key
user_id : '1'


```
create table users (
    user_id int,
    name text,
    age int,
    gender boolean,
    PRIMARY KEY(user_id))


SELECT name FROM users WHERE user_id=1



SELECT * FROM users WHERE name="Thomas"

SELECT * FROM users WHERE age>45 and name="Thomas"
```

Not possible out of the box without secondary Index or further extra tables

No problem with DSE Search

What is the value of DSE Search?

Multi-criteria WHERE Constraints

- WHERE constraints with multiple columns

- No extra tables needed

Full Text Search

- **Wildcards** ? *, like or Lemmatisation

- **Faceting**, Slice and Dice

Suggester

- Implement type-ahead functionality with **Fuzzy** or **AnalyticInFix** Lookup Factories

Geospatial queries

- Queries with coordinates and distances search

DATASTAX

Integrated with CQL

1. Create core, schema.xml and solrconfig.xml.

2. [optional] customize schema.xml, solrconfig.xml

3. Start indexing, re-indexing

```
Wildcard search
cqlsh> select * from sales where solr_cql = ' {`fq`:`name=gre*`} '


Facet query
cqlsh> SELECT * FROM sales WHERE solr_query='{"q":"name:*","facet":{"field":"item"}}';


Range query
Cqlsh> SELECT * FROM sales WHERE solr_query='{"q":"dt:[2017-01-01 TO 2017-01-10]"}';


Geo search
cqlsh> select * from sales where solr_cql =
                    ' {"q":"*:*", "fq":"point:\"IsWithin(BUFFER(POINT(4.0 49.0), 10.0))\""}'
```

DSE Search Architecture

JVM

security Monitoring CQL

C* field type support (UDT, Tuple, collections)
automatic index updates
live indexing
fault-tolerant distributed search

multi-DC replication partitioning repair **cassandra**

Solr — geospatial text analysis caching

lucene — sorting wildcards grouping faceting

With all benefits:

High Available

Scalable

Multi Data Center

Index Security

Integrated in the the same JVM

- Data Locality (TokenRanges) and Shared Memory
- High Available, no master needed
- Index is sharded to all nodes

Transparency

- Access via CQL or REST API

## Automated indexing with INSERT / UPDATE

- No extra ETL Process needed

The Open Source Way



Application

DB API | Solr API

Manuel ETL

Re-Indexing …

Apache Cassandra™
Cluster

Apache Solr™
Cluster

## Separated Search Cluster

- "Split Brain" risk, data inconsistency
- ETL to generate, update and re-create index

Complex application

Two seperated APIs, driver, read paths

https://docs.datastax.com/en/datastax_enterprise/5.0/datastax_enterprise/srch/searchOssSolrDiff.html

Inserting through CQL

Querying through CQL

```
SELECT title FROM solr WHERE solr_query='title:natio*';


 title
-------------------------------------------------------------------
Bolivia national football team 2002
List of French born footballers who have played for other national teams
Lithuania national basketball team at Eurobasket 2009
Bolivia national football team 2000
Kenya national under-20 football team
Bolivia national football team 1999
Israel men's national inline hockey team
Bolivia national football team 2001
```
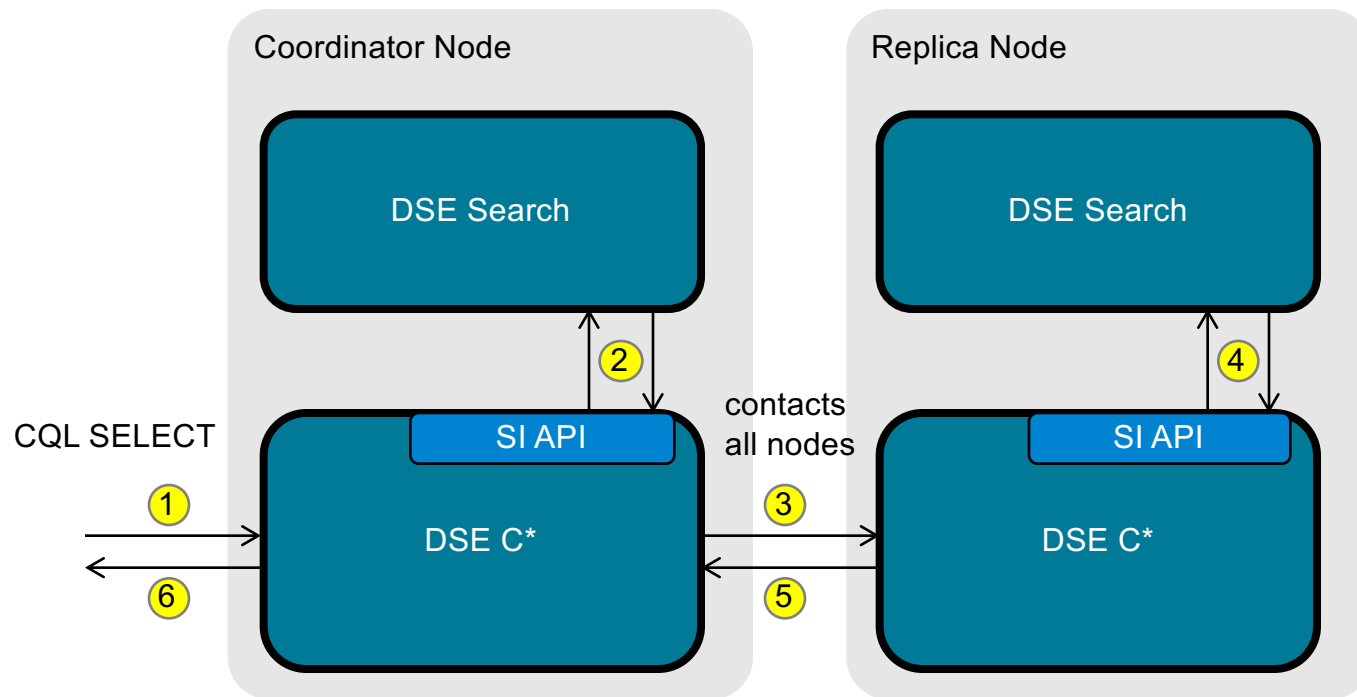
How many nodes to contact?
- We don't know the primary key
- Theory:
    - contact at least one replica for every token range

- Cassandra contacts all nodes
- **Our custom Solr SearchComponent does intelligent shard selection**

Querying through CQL

Lab 5 : Hands-on DSE Search

The power behind the moment. | DATASTAX

Thank You!